

Exploring data 2

Other R objects: Matrices and lists

Matrices

A matrix is like a data frame, but all the values in all columns must be of the same class (e.g., numeric, character). (Another way you can think of it is as a “wrapped” vector.)

Matrices can be faster and more memory-efficient than data frames. Also, a lot of statistical methods within R code is implemented using linear algebra and other mathematical techniques based on matrices.

Matrices

We can use the `matrix()` function to construct a matrix:

```
foo <- matrix(1:10, ncol = 5)
foo
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    3    5    7    9
## [2,]    2    4    6    8   10
```

Matrices

The `as.matrix()` function is used to convert an object to a matrix:

```
foo <- data.frame(col_1 = 1:2, col_2 = 3:4,  
                  col_3 = 5:6, col_4 = 7:8,  
                  col_5 = 9:10)  
foo <- as.matrix(foo)  
foo
```

```
##      col_1 col_2 col_3 col_4 col_5  
## [1,]     1     3     5     7     9  
## [2,]     2     4     6     8    10
```

Matrices

You can index matrices with square brackets, just like data frames:

```
foo[1, 1:2]
```

```
## col_1 col_2  
##      1      3
```

You cannot, however, use dplyr functions with matrices:

```
foo %>% filter(col_1 == 1)
```

```
Error in UseMethod("filter_") :  
  no applicable method for 'filter_' applied to an object of  
  class "c('matrix', 'integer', 'numeric')"
```

Lists

Lists are the “kitchen sink” of R objects. They can be used to keep together a variety of different R objects of different classes, dimensions, and structures in a single R object.

Because there are often cases where an R operation results in output that doesn't have a simple structure, lists can be a very useful way to output complex output from an R function.

Most lists are not “tidy” data. However, we'll cover some ways that you can easily “tidy” some common list objects you might use a lot in your R code, including the output of fitting linear and generalized linear models.

Lists

```
example_list <- list(a = sample(1:10, 5),  
                    b = tibble(letters = letters[1:3],  
                               numbers = 1:3))
```

```
example_list
```

```
## $a  
## [1] 9 7 3 1 5  
##  
## $b  
## # A tibble: 3 x 2  
##   letters numbers  
##   <chr>      <int>  
## 1 a          1  
## 2 b          2  
## 3 c          3
```

Indexing lists

To pull an element out of a list, you can either use \$ or [[]] indexing:

```
example_list$a
```

```
## [1] 9 7 3 1 5
```

```
example_list[[2]]
```

```
## # A tibble: 3 x 2
```

```
##   letters numbers
```

```
##   <chr>      <int>
```

```
## 1 a          1
```

```
## 2 b          2
```

```
## 3 c          3
```

Indexing lists

To access a specific value within a list element we can index the element using double, double brackets:

```
example_list[["b"]][["numbers"]]
```

```
## [1] 1 2 3
```

Again, we can index using names or numeric indices:

```
example_list[["b"]][[1]]
```

```
## [1] "a" "b" "c"
```

Exploring lists

If an R object is a list, running `class` on the object will return “list”:

```
class(example_list)
```

```
## [1] "list"
```

Often, lists will have names for each element (similar to column names for a dataframe). You can get the names of all elements of a list using the `names` function:

```
names(example_list)
```

```
## [1] "a" "b"
```

Exploring lists

The `str` function is also useful for exploring the structure of a list object:

```
str(example_list)
```

```
## List of 2
## $ a: int [1:5] 9 7 3 1 5
## $ b:Classes 'tbl_df', 'tbl' and 'data.frame':  3 obs. of  2 variables:
## ..$ letters: chr [1:3] "a" "b" "c"
## ..$ numbers: int [1:3] 1 2 3
```

Exploring lists

A list can even contain other lists. We can use the `str` function to see the structure of a list:

```
a_list <- list(list("a", "b"), list(1, 2))
```

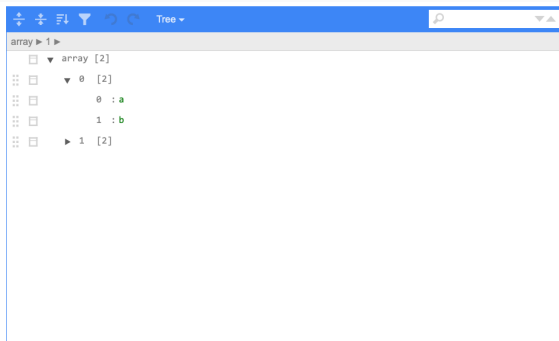
```
str(a_list)
```

```
## List of 2
## $ :List of 2
## ..$ : chr "a"
## ..$ : chr "b"
## $ :List of 2
## ..$ : num 1
## ..$ : num 2
```

Exploring lists

Using `str` to print out the list's structure doesn't produce the easiest to digest output. We can use the `jsonedit` function from the `listviewer` package to create a widget in the Viewer pane to more easily explore our list.

```
library(listviewer)
jsonedit(a_list)
```



Lists versus dataframes

As a note, a dataframe is actually just a very special type of list. It is a list where every element (column in the dataframe) is a vector of the same length, and the object has a special attribute specifying that it is a dataframe.

```
example_df <- tibble(letters = letters[1:3],  
                     number = 1:3)
```

```
class(example_df)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

```
is.list(example_df)
```

```
## [1] TRUE
```


We'll take a break now to do section 1 of the In-course Exercise for Chapter 7.

Simple statistical tests in R

R has many different functions for different statistical tests.

So far, I have not found a statistical test I wanted to perform that did not have a function to do it in R.

We'll start with a pretty simple test called the “Shapiro-Wilk test of normality.”

The **null hypothesis** for this test is that the data follow a normal distribution. If the p-value from running the test is lower than a specified threshold (often 0.05), then you reject the null hypothesis.

Often, a parametric test will require the assumption that one or more variables follow a normal distribution. This test can help check that assumption.

If you find that the distribution does not seem to be normal, you can make adjustments, like transforming the variable or using a non-parametric test (see [here](#) for a nice overview).

Statistical tests

We'll use this test as an example, but it has two characteristics that are common to most simple statistical tests functions in R:

1. It inputs a vector (or vectors, for some tests), rather than a tidy dataframe.
2. It outputs a list object, rather than a dataframe.

We'll talk about both of these characteristics and how we can still work statistical tests into a “tidy” workflow, where we're keeping the data in a tidy dataframe format most of the time.

Statistical tests

To show how this test works, let's simulate some data that we know are normal.

There are several functions in R that let you simulate a vector of data from a specified distribution. For a normal distribution, the function is `rnorm`.

For this function, you'll specify how many values you want in the vector (`n`) and the values of the mean (`mean`) and standard deviation (`sd`) of the normal distribution.

Statistical tests

For example, you can run this to create a vector with a random sample of 1,000 values from a normal distribution with mean 200 and standard deviation 50:


```
normal_ex_vector <- rnorm(n = 1000, mean = 200, sd = 50)
head(normal_ex_vector)
```

```
## [1] 269.0354 171.9940 208.1412 230.6753 189.9154 128.6256
```


Statistical tests

Let's check the histogram of this data. Since this data is in a vector, not a dataframe, we should use `qplot` from `ggplot2` instead of `ggplot`:

```
qplot(normal_ex_vector, geom = "histogram")
```



Statistical tests

It looks pretty normal, but let's run the test. The function to run the test is `shapiro.test`. Its only argument is `x`, the vector that you want to check.

```
shapiro.test(x = normal_ex_vector)
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  normal_ex_vector  
## W = 0.9987, p-value = 0.6919
```

The default for this function is to print out the results in a way that's easy to read on the screen. However, you can also save the results from the test in a new object:

```
ex_sw_result <- shapiro.test(x = normal_ex_vector)
```

This object has a special class called “htest”:

```
class(ex_sw_result)
```

```
## [1] "htest"
```

Really, though, this is just a special kind of list:

```
is.list(ex_sw_result)
```

```
## [1] TRUE
```

Statistical tests

If you run `str` on the object, you can see it has all the information that is usually printed by the function tucked away in different slots of the list:

```
str(ex_sw_result)
```

```
## List of 4
## $ statistic: Named num 0.999
##   ..- attr(*, "names")= chr "W"
## $ p.value   : num 0.692
## $ method    : chr "Shapiro-Wilk normality test"
## $ data.name: chr "normal_ex_vector"
## - attr(*, "class")= chr "htest"
```

Statistical tests

There is a package called `broom` that can pull this information out and format it as a tidy dataframe. For example:

```
library(broom)
tidy(ex_sw_result)
```

```
## # A tibble: 1 x 3
##   statistic p.value method
##   <dbl>    <dbl> <chr>
## 1      0.999    0.692 Shapiro-Wilk normality test
```

This seems unexciting for this example, but trust me, it turns out that being able to do this is **very** exciting.

Statistical tests

Now, let's look at how this process would work if you were starting with a vector that was a column in a tidy dataframe. Since you're using tidyverse tools, you'll probably find you want to do this often.

We'll create a very simple dataframe with only this column:


```
ex_df <- tibble(norm_vector = normal_ex_vector)
ex_df %>%
  slice(1:3)
```

```
## # A tibble: 3 x 1
##   norm_vector
##         <dbl>
## 1         269.
## 2         172.
## 3         208.
```

Statistical tests

Now you can use `ggplot` to make the histogram:

```
ggplot(ex_df, aes(x = norm_vector)) +  
  geom_histogram()
```



Statistical tests

To fit the test, you'll need to be able to pull this vector out of the dataframe. To do that, you can use the `pluck` function from the `purrr` package in a pipeline. That function “plucks” out a single column as a vector. For example:

```
library(purrr)
ex_df %>%
  pluck("norm_vector") %>%
  head()

## [1] 269.0354 171.9940 208.1412 230.6753 189.9154 128.6256
```

Statistical tests

With that function, you can pipe right into the Shapiro-Wilk test function:

```
ex_df %>%  
  pluck("norm_vector") %>%  
  shapiro.test()  
  
##  
##  Shapiro-Wilk normality test  
##  
## data:  .  
## W = 0.9987, p-value = 0.6919
```

Statistical tests

Now just add on the `tidy` function to get the test output in a tidy dataframe, and you're back to your typical format!

```
ex_df %>%  
  pluck("norm_vector") %>%  
  shapiro.test() %>%  
  tidy()
```

```
## # A tibble: 1 x 3  
##   statistic p.value method  
##   <dbl>    <dbl> <chr>  
## 1      0.999    0.692 Shapiro-Wilk normality test
```

Now let's look at some real data. The variable dataframe of the atlas1006 dataset in the microbiome library has a column on diversity.

We might want to test if diversity is different by gender, nationality, or other factors. To pick which statistical tests to use to check those questions, though, it will help to know if this variable is normally distributed.

Statistical tests

The atlas1006 data is stored in a phyloseq object (think of it as a fancy type of list). To extract a dataframe with characteristics of the samples, you'll need to use `get_variable` (which we can pipe into if we want):

```
library(microbiome)
data(atlas1006)
atlas1006 %>%
  get_variable() %>%
  slice(1:3)
```

##	age	sex	nationality	DNA_extraction_method	project	divers
## 1	28	male	US	<NA>	1	5
## 2	24	female	US	<NA>	1	6
## 3	52	male	US	<NA>	1	5

##	bmi_group	subject	time	sample
## 1	severeobese	1	0	Sample-1
## 2	obese	2	0	Sample-2
## 3	lean	3	0	Sample-3

There are a few people that they measure several times, so there are more rows than the number of people they measure:

```
atlas1006 %>%  
  get_variable() %>%  
  nrow()
```

```
## [1] 1151
```

Statistical tests

We probably just want to work with the first measurement from each person, so let's use `filter` to filter to samples with a “time” value of 0 (first measurement):

```
atlas1006 %>%  
  get_variable() %>%  
  filter(time == 0) %>%  
  nrow()
```


```
## [1] 1006
```

This looks right.

Statistical tests

We can use a histogram to visually check the normality:

```
atlas1006 %>%  
  get_variable() %>%  
  filter(time == 0) %>%  
  ggplot(aes(x = diversity)) +  
  geom_histogram()
```



To extract the column on diversity as a vector, we can use pluck:

```
atlas1006 %>%  
  get_variable() %>%  
  filter(time == 0) %>%  
  pluck("diversity") %>%  
  head()
```

```
## [1] 5.76 6.06 5.50 5.87 5.89 5.53
```

Statistical tests

Now add on the Shapiro test function:

```
atlas1006 %>%  
  get_variable() %>%  
  filter(time == 0) %>%  
  pluck("diversity") %>%  
  shapiro.test()  
  
##  
##  Shapiro-Wilk normality test  
##  
## data:  .  
## W = 0.93439, p-value < 2.2e-16
```

Statistical tests

And finally add on the tidy function:

```
atlas1006 %>%  
  get_variable() %>%  
  filter(time == 0) %>%  
  pluck("diversity") %>%  
  shapiro.test() %>%  
  tidy()
```

```
## # A tibble: 1 x 3  
##   statistic p.value method  
##   <dbl>    <dbl> <chr>  
## 1      0.934 1.29e-20 Shapiro-Wilk normality test
```

Find out more about statistical tests in R

I won't be teaching in this course how to find the correct statistical test. That's something you'll hopefully learn in a statistics course.

There are also a variety of books that can help you with this, including some that you can access free online through CSU's library. One servicable introduction is "Statistical Analysis with R for Dummies".

We'll take a break now to do sections 2 and 3 of the In-course Exercise for Chapter 7.

Regression models

World Cup example

In the World Cup data, we may wonder if the number of tackles is associated with the time the player played. Let's start by grabbing just the variables we care about (we'll be using Position later, so we'll include that):


```
library(faraway)
data(worldcup)
worldcup <- worldcup %>%
  select(Time, Tackles, Position)
worldcup %>% slice(1:3)
```

##	Time	Tackles	Position
## 1	16	0	Midfielder
## 2	351	14	Midfielder
## 3	180	6	Defender

World Cup example

We can start by plotting the relationship between the time a player played and the number of tackles they had:

```
library(ggplot2)
ggplot(worldcup, aes(Time, Tackles)) +
  geom_point()
```



World Cup example

There does indeed seem to be an association. Next, we might want to test this using some kind of statistical model or test.

Let's start by fitting a linear regression model, to see if there's evidence that tackles tend to change (increase or decrease) as the player's time played increases.

(In a bit, we'll figure out that a linear model might not be the best way to model this, since the number of tackles is a count, rather than a variable with a normal distribution, but bear with me. . .)

Formula structure

Regression models can be used to estimate how the expected value of a *dependent variable* changes as *independent variables* change.

In R, regression formulas take this structure:

Generic code

```
[response variable] ~ [indep. var. 1] + [indep. var. 2] + ...
```

Notice that ~ used to separate the independent and dependent variables and the + used to join independent variables. This format mimics the statistical notation:

$$Y_i \sim X_1 + X_2 + \cdots + \epsilon_i$$

You will use this type of structure in R for a lot of different function calls, including those for linear models (`lm`) and generalized linear models (`glm`).

Linear models

To fit a linear model, you can use the function `lm()`. Use the `data` option to specify the dataframe from which to get the vectors. You can save the model as an object.

```
tackle_model <- lm(Tackles ~ Time, data = worldcup)
```

This call fits the model:

$$Y_i = \beta_0 + \beta_1 X_{1,i} + \epsilon_i$$

where:

- Y_i : Number of tackles for player i (dependent variable)
- $X_{1,i}$: Minutes played by player i (independent variable)

A few things to point out:

- By default, an intercept is fit to the model.
- If you specify a dataframe using `data` in the `lm` call, you can write the model formula using just the column names for the independent variable(s) and dependent variable you want, without quotation marks around those names.
- You can save the output of fitting the model to an R object (if you don't, a summary of the fit model will be print out at the console).

Model objects

The output from fitting a model using `lm` is a list object:

```
class(tackle_model)
```

```
## [1] "lm"
```

This list object has a lot of different information from the model, including overall model summaries, estimated coefficients, fitted values, residuals, etc.

```
names(tackle_model)
```

```
## [1] "coefficients" "residuals"      "effects"         "rank"
## [5] "fitted.values" "assign"          "qr"              "df.residual"
## [9] "xlevels"       "call"           "terms"           "model"
```

Model objects and broom

This list object is not in a “tidy” format. However, you can use functions from broom to pull “tidy” dataframes from this model object.

For example, you can use the glance function to pull out a one-row tidy dataframe with model summaries.

```
glance(tackle_model)
```

```
## # A tibble: 1 x 11
##   r.squared adj.r.squared sigma statistic p.value    df logLik   AIC   BIC
##   <dbl>      <dbl> <dbl>    <dbl>    <dbl> <int>  <dbl> <dbl> <dbl>
## 1    0.373      0.372  3.69    353. 4.31e-62     2 -1620. 3246. 3259.
## # ... with 2 more variables: deviance <dbl>, df.residual <int>
```

Model objects and broom

If you want to get the estimated model coefficients (and some related summaries) instead, you can use the `tidy` function to do that:

```
tidy(tackle_model)
```

```
## # A tibble: 2 x 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)    0.110    0.265     0.415 6.78e- 1
## 2 Time           0.0195   0.00104   18.8  4.31e-62
```

This output includes, for each model term, the **estimated coefficient** (estimate), its **standard error** (std.error), the **test statistic** (for `lm` output, the statistic for a test with the null hypothesis that the model coefficient is zero), and the associated **p-value** for that test (p.value).

Model objects and broom

Some of the model output have a value for each original observation (e.g., fitted values, residuals). You can use the `augment` function to add those elements to the original data used to fit the model:


```
augment(tackle_model) %>%  
  slice(1:2)
```

```
## # A tibble: 2 x 10  
##   .rownames Tackles   Time .fitted .se.fit .resid   .hat .sigma .cooksd  
##   <chr>      <int> <int>   <dbl>   <dbl> <dbl>   <dbl> <dbl>   <dbl>  
## 1 Abdoun         0    16   0.423   0.251 -0.423  0.00464  3.69 3.07e-5  
## 2 Abe           14   351   6.97    0.212  7.03  0.00329  3.68 6.01e-3  
## # ... with 1 more variable: .std.resid <dbl>
```


Model objects and broom

One important use of this augment output is to create a plot with both the original data and a line showing the fit model (via the predictions):

```
augment(tackle_model) %>%  
  ggplot(aes(x = Time, y = Tackles)) +  
  geom_point(size = 0.8, alpha = 0.8) +  
  geom_line(aes(y = .fitted), color = "red", size = 1.2)
```




Model objects and autoplot

There is a function called `autoplot` in the `ggplot2` package that will check the class of an object and then create a certain default plot for that class. Although the generic `autoplot` function is in the `ggplot2` package, for `lm` and `glm` objects, you must have the `ggfortify` package installed and loaded to be able to access the methods of `autoplot` specifically for these object types.

If you have the package that includes an `autoplot` method for a specific object type, you can just run `autoplot` on the objects name and get a plot that is considered a useful default for that object type. For `lm` objects, `autoplot` gives small graphics with model diagnostic plots.

Model objects and autoplot

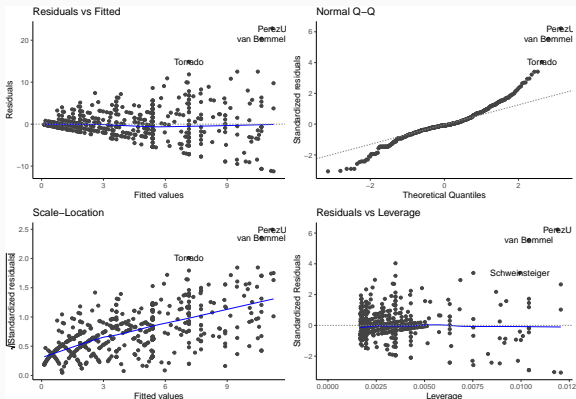
```
library(ggfortify)
autoplot(tackle_model)
```



Model objects and autoplot

The output from autoplot is a ggplot object, so you can add elements to it as you would with other ggplot objects:

```
autoplot(tackle_model) +  
  theme_classic()
```

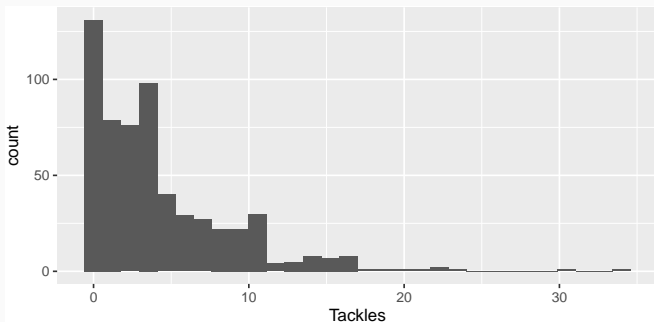


Regression models

In this case, these diagnostics clearly show that there are some problems with using a linear regression model to fit this data.

Many of these issues arise because the outcome (dependent) variable doesn't follow a normal distribution.

```
ggplot(worldcup, aes(x = Tackles)) +  
  geom_histogram()
```



A better model, therefore, might be one where we assume that Tackles follows a Poisson distribution, rather than a normal distribution. (For variables that represent counts, this will often be the case.)

In the a little bit, we'll look at **generalized linear models**, which let us extend the idea of a linear model to situations where the dependent variable follows a distribution other than the normal distribution.

We'll take a break now to do section 4 of the In-Course Exercise for Chapter 7.

Fitting a model with a factor

You can also use binary variables or factors (i.e., categorical variables) as independent variables in regression models:

```
tackles_model_2 <- lm(Tackles ~ Position, data = worldcup)
```

This call fits the model:

$$Y_i = \beta_0 + \beta_1 X_{1,i} + \epsilon_i$$

where $X_{1,i}$: Position of player i

Fitting a model with a factor

If there are more than one levels to the factor, then the model will fit a separate value for each level of the factor above the first level (which will serve as a baseline):

```
levels(worldcup$Position)
```

```
## [1] "Defender"    "Forward"     "Goalkeeper"  "Midfielder"
```

```
tidy(tackles_model_2)
```

```
## # A tibble: 4 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	5.46	0.315	17.3	1.14e-54
## 2	PositionForward	-3.44	0.480	-7.17	2.20e-12
## 3	PositionGoalkeeper	-5.43	0.787	-6.91	1.27e-11
## 4	PositionMidfielder	-0.300	0.426	-0.705	4.81e- 1

Fitting a model with a factor

The intercept is the expected (average) value of the outcome (Tackles) for the first level of the factor. Each other estimate gives the expected difference between the value of the outcome for this first level of Position and one of the other levels of the factor.

```
levels(worldcup$Position)
```

```
## [1] "Defender"    "Forward"     "Goalkeeper"  "Midfielder"
```

```
tidy(tackles_model_2)
```

```
## # A tibble: 4 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	5.46	0.315	17.3	1.14e-54
## 2	PositionForward	-3.44	0.480	-7.17	2.20e-12
## 3	PositionGoalkeeper	-5.43	0.787	-6.91	1.27e-11
## 4	PositionMidfielder	-0.300	0.426	-0.705	4.81e- 1

Linear models versus GLMs

You can fit a variety of models, including linear models, logistic models, and Poisson models, using generalized linear models (GLMs).

For linear models, the only difference between `lm` and `glm` is how they're fitting the model (least squares versus maximum likelihood). You should get the same results regardless of which you pick.

Linear models versus GLMs

For example:

```
glm(Tackles ~ Time, data = worldcup) %>%  
  tidy()
```

```
## # A tibble: 2 x 5  
##   term          estimate std.error statistic  p.value  
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept)    0.110    0.265     0.415 6.78e- 1  
## 2 Time          0.0195   0.00104   18.8  4.31e-62
```

```
lm(Tackles ~ Time, data = worldcup) %>%  
  tidy()
```

```
## # A tibble: 2 x 5  
##   term          estimate std.error statistic  p.value  
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept)    0.110    0.265     0.415 6.78e- 1  
## 2 Time          0.0195   0.00104   18.8  4.31e-62
```

You can fit other model types with `glm()` using the `family` option:

Model type	family option
Linear	<code>family = gaussian(link = 'identity')</code>
Logistic	<code>family = binomial(link = 'logit')</code>
Poisson	<code>family = poisson(link = 'log')</code>

GLM example

For example, say we wanted to fit a GLM, but specifying a Poisson distribution for the outcome (and a log link) since we think that Tackles might be distributed with a Poisson distribution:


```
tackle_model_3 <- glm(Tackles ~ Time, data = worldcup,  
                      family = poisson(link = "log"))  
  
tackle_model_3 %>%  
  tidy()
```

```
## # A tibble: 2 x 5  
##   term          estimate std.error statistic  p.value  
##   <chr>          <dbl>     <dbl>     <dbl>    <dbl>  
## 1 (Intercept)  0.350     0.0443      7.90 2.85e- 15  
## 2 Time        0.00422   0.000129    32.8 1.81e-235
```

GLM example

Here are the predicted values from this model (red line):

```
tackle_model_3 %>%  
  augment() %>%  
  mutate(.fitted = exp(.fitted)) %>%  
  ggplot(aes(x = Time, y = Tackles)) +  
  geom_point() +  
  geom_line(aes(y = .fitted), color = "red", size = 1.2)
```



Formula structure

There are some conventions that can be used in R formulas. Common ones include:

Convention	Meaning
I()	calculate the value inside before fitting (e.g., I(x1 + x2))
:	fit the interaction between two variables (e.g., x1:x2)
*	fit the main effects and interaction for both variables (e.g., x1*x2 equals x1 + x2 + x1:x2)
.	fit all variables other than the response (e.g., y ~ .)
-	do not include a variable (e.g., y ~ . - x1)
1	intercept (e.g., y ~ 1)

To find out more

Great resources to find out more about using R for basic statistics:

- Statistical Analysis with R for Dummies, Joseph Schmuller (free online through our library; Chapter 14 covers regression modeling)
- The R Book, Michael J. Crawley (free online through our library; Chapter 14 covers regression modeling, Chapters 10 and 13 cover linear and generalized linear regression modeling)
- R for Data Science (Section 4)

If you want all the details about fitting linear models and GLMs in R, Faraway's books are fantastic (more at level of Master's in Applied Statistics):

- Linear Models with R, Julian Faraway (also freely available online through our library)
- Extending the Linear Model with R, Julian Faraway (available in hardcopy through our library)

We'll take a break now to do section 5 of the In-Course Exercise for Chapter 7.